



Moonshot: Implementing the Next Generation of Network Telemetry Technologies

Andy Gospodarek
Broadcom Corporation
gospo@broadcom.com

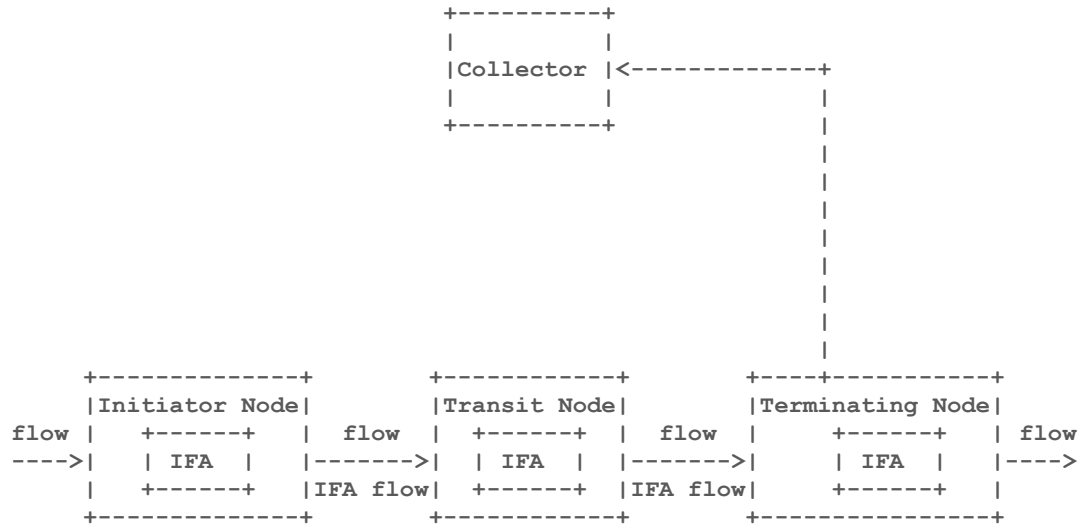
Introduction to Network Telemetry



Introduction to Network Telemetry

- Generic network-wide method for reporting and collecting network state information on packets as they traverse the network
- Packets contain *telemetry instructions* and each supported device responds to those instructions by adding metadata to frames
- *Telemetry instructions* can be added *In-Situ* to existing frames or to sampled frames that are transmitted simultaneously with existing traffic
- Standards like IOAM, INT, and IFA all describe ways to implement Network Telemetry

Network Telemetry Components





Source or Initiator Node

```
+-----+
|Initiator Node|
flow | +-----+ | flow
---->| | IFA | |----->
    | +-----+ | IFA flow
+-----+
```

Node that adds telemetry instructions to frames -- in-band or out-of-band



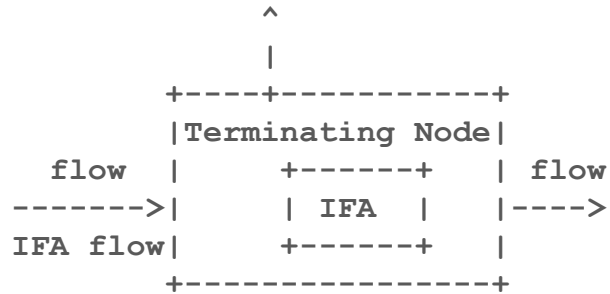
Transit Hop or Transit Node

```
      +-----+
      |Transit Node|
flow  | +-----+ | flow
----->| | IFA | |----->
IFA flow| +-----+ | IFA flow
      +-----+
```

Any node that adds telemetry metadata to frame that contains telemetry instructions




Sink or Terminating Node



Node that removes telemetry headers from frames and sends them to a collector



Collector



Any application that will receive and process telemetry data collected by Sink or Terminating Node

Inband Network Telemetry (INT)




Inband Network Telemetry (INT)

- Framework suggested by P4.org
- True In-band monitoring as frames are modified as they traverse the network
- Compatible with encapsulation formats that have extension/option capabilities
 - INT over VXLAN (as VXLAN payload, per GPE extension)
 - INT over Geneve (as Geneve option)
 - INT over GRE (as a shim between GRE header and encapsulated payload)
 - INT over NSH (as NSH payload)
- Support for IP/IPv6 TCP/UDP traffic not as easy since there is not room for extensions
 - INT header and metadata transparently added to payload



INT IPv4/TCP Frame Format



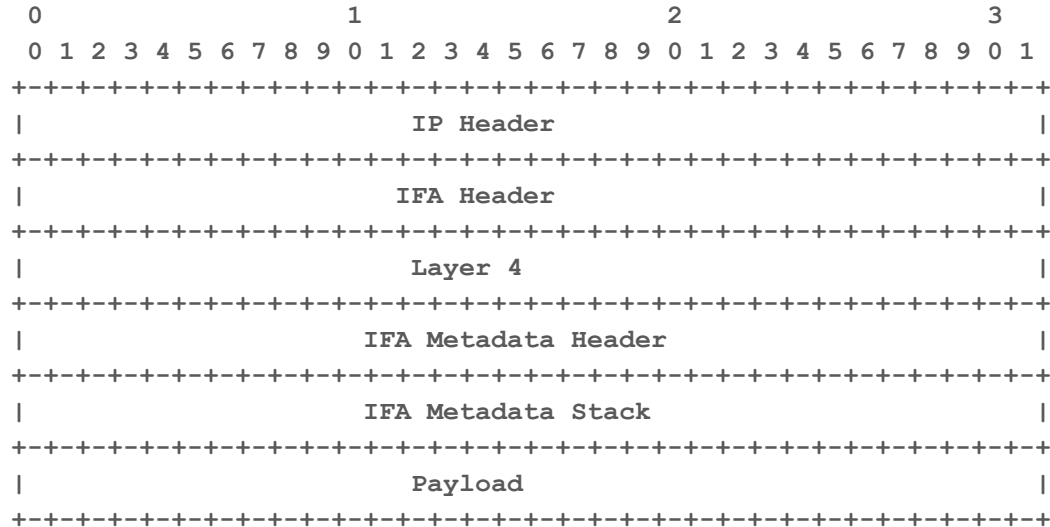


Inband Flow Analyzer (IFA)

- Specification released later to address perceived shortcomings of IOAM and INT
- Adds ability to sent telemetry instructions in sampled/copied frames rather than original frames
- Frame format aims to be compatible with more hardware

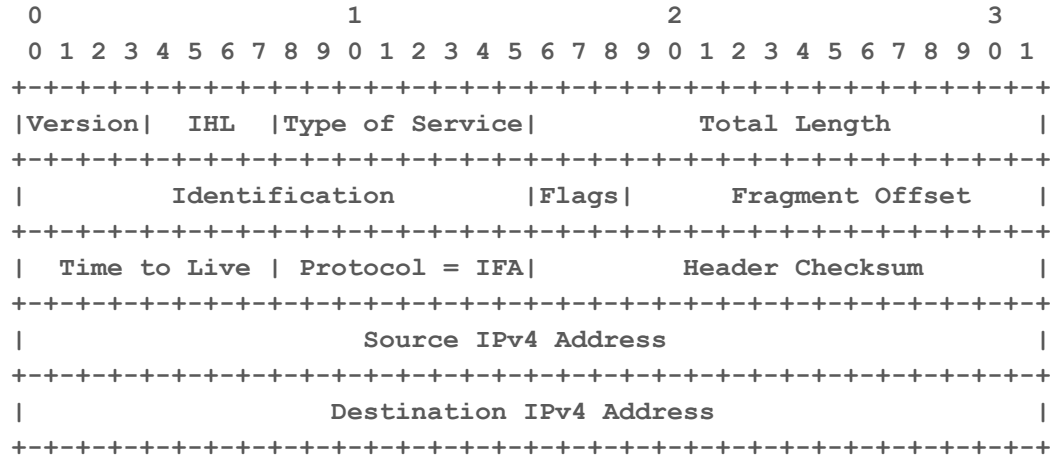


IFA IPv4/TCP Frame Format





IFA IPv4 Header



Telemetry Software Architecture



Source or Initiator Node

Desired Functionality	Kernel Implementation
Packet Sampling	TC or netfilter
Packet Mirroring (IFA)	TC
Packet Redirect	TC or XDP/eBPF
Encapsulate and Transmit	new lwtnet type or XDP/eBPF



Transit Hop or Transit Node

Desired Functionality	Kernel Implementation
Match on INT/IFA Frame	TC
Update Metadata and Transmit	new lwtnnelt type or XDP/eBPF



Sink or Terminating Node

Desired Functionality	Kernel Implementation
Match on INT/IFA Frame	TC or netfilter
Collect metadata from frame	new lwtunnel type or XDP/eBPF
Send frame to collector	new lwtunnel type or XDP/eBPF
Transmit original frame (INT)	new lwtunnel type or XDP/eBPF

Thinking Beyond the Linux Kernel Datapath



Hardware Support

- Switch Hardware
 - INT is by P4
 - Switch vendors and switch ASIC vendors are active in INT and IFA standards writing
- NIC Hardware
 - NIC vendors are active in INT and IFA standards writing
 - NIC+FPGA solutions for IOAM exist -- ASICs will add support soon



Network-based configuration

- Kernel datapath examples all presume configuration by netlink
- As hardware support for IFA is added, expect configuration over network
 - Common in switches
 - Becoming more common in servers used for baremetal deployments



Userspace Implementations

- Unlikely to be a DPDK application just for INT/IFA
- Libraries for DPDK-based applications could be created to add INT/IFA functionality to existing dataplane applications

Development and Deployment Risks



User Risk

- Protocols and packet types can change
- Risk is minimized when single organization controls entire infrastructure



Community Risk

- Early acceptance to Linux kernel could be problematic if software not kept up to date across infrastructure
- Risk is minimized when single organization controls entire infrastructure

Děkuji!